

Assignment_3

October 6, 2022

Harshath.M
732119104028

#Unzipping

```
[ ]: #!unzip '/content/drive/MyDrive/Colab Notebooks/Flowers-Dataset.zip'
```

#Data Augmentation

```
[ ]: from tensorflow.keras.preprocessing.image import ImageDataGenerator

train_gen = ImageDataGenerator(rescale=1./255,
                                zoom_range=0.2,
                                horizontal_flip=True)
test_gen = ImageDataGenerator(rescale=1./255)
```

```
[ ]: xtrain = train_gen.flow_from_directory('/content/flowers',
                                           target_size=(64,64),
                                           class_mode='categorical',
                                           batch_size=100)
```

Found 4317 images belonging to 5 classes.

#Train

```
[ ]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Convolution2D, MaxPooling2D, Dense, Flatten
from keras.callbacks import EarlyStopping, ReduceLROnPlateau
```

```
[ ]: model = Sequential()
model.add(Convolution2D(32,(3,3),activation='relu',input_shape=(64,64,3)))
model.add(MaxPooling2D((2,2)))
model.add(Flatten())

model.add(Dense(400,activation='relu'))
model.add(Dense(200,activation='relu'))
model.add(Dense(100,activation='relu'))
model.add(Dense(5,activation='softmax'))
```

```
[ ]: model.  
      ↪ compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
[ ]: early_stopping = EarlyStopping(monitor='accuracy',  
                                   patience=3)  
      reduce_lr = ReduceLROnPlateau(monitor='accuracy',  
                                   patience=5,  
                                   factor=0.5, min_lr=0.00001)  
  
      callback = [reduce_lr, early_stopping]
```

```
[ ]: model.fit_generator(xtrain,  
                        steps_per_epoch = len(xtrain),  
                        callbacks=callback,  
                        epochs=100)
```

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:4: UserWarning:
`Model.fit_generator` is deprecated and will be removed in a future version.
Please use `Model.fit`, which supports generators.
after removing the cwd from sys.path.

```
Epoch 1/100  
44/44 [=====] - 31s 703ms/step - loss: 0.6074 -  
accuracy: 0.7751 - lr: 0.0010  
Epoch 2/100  
44/44 [=====] - 31s 702ms/step - loss: 0.5491 -  
accuracy: 0.7973 - lr: 0.0010  
Epoch 3/100  
44/44 [=====] - 31s 696ms/step - loss: 0.5417 -  
accuracy: 0.8043 - lr: 0.0010  
Epoch 4/100  
44/44 [=====] - 31s 693ms/step - loss: 0.4930 -  
accuracy: 0.8156 - lr: 0.0010  
Epoch 5/100  
44/44 [=====] - 31s 692ms/step - loss: 0.4616 -  
accuracy: 0.8293 - lr: 0.0010  
Epoch 6/100  
44/44 [=====] - 31s 695ms/step - loss: 0.4350 -  
accuracy: 0.8392 - lr: 0.0010  
Epoch 7/100  
44/44 [=====] - 31s 700ms/step - loss: 0.4190 -  
accuracy: 0.8469 - lr: 0.0010  
Epoch 8/100  
44/44 [=====] - 31s 692ms/step - loss: 0.3975 -  
accuracy: 0.8568 - lr: 0.0010  
Epoch 9/100  
44/44 [=====] - 31s 689ms/step - loss: 0.4207 -  
accuracy: 0.8432 - lr: 0.0010
```

Epoch 10/100
44/44 [=====] - 31s 696ms/step - loss: 0.3674 -
accuracy: 0.8687 - lr: 0.0010
Epoch 11/100
44/44 [=====] - 31s 689ms/step - loss: 0.3267 -
accuracy: 0.8888 - lr: 0.0010
Epoch 12/100
44/44 [=====] - 31s 699ms/step - loss: 0.3255 -
accuracy: 0.8877 - lr: 0.0010
Epoch 13/100
44/44 [=====] - 31s 688ms/step - loss: 0.2975 -
accuracy: 0.8999 - lr: 0.0010
Epoch 14/100
44/44 [=====] - 31s 690ms/step - loss: 0.3144 -
accuracy: 0.8927 - lr: 0.0010
Epoch 15/100
44/44 [=====] - 31s 697ms/step - loss: 0.2544 -
accuracy: 0.9120 - lr: 0.0010
Epoch 16/100
44/44 [=====] - 31s 694ms/step - loss: 0.2986 -
accuracy: 0.8939 - lr: 0.0010
Epoch 17/100
44/44 [=====] - 31s 699ms/step - loss: 0.2570 -
accuracy: 0.9097 - lr: 0.0010
Epoch 18/100
44/44 [=====] - 31s 691ms/step - loss: 0.2282 -
accuracy: 0.9205 - lr: 0.0010
Epoch 19/100
44/44 [=====] - 31s 697ms/step - loss: 0.2180 -
accuracy: 0.9226 - lr: 0.0010
Epoch 20/100
44/44 [=====] - 31s 694ms/step - loss: 0.1995 -
accuracy: 0.9314 - lr: 0.0010
Epoch 21/100
44/44 [=====] - 31s 690ms/step - loss: 0.1805 -
accuracy: 0.9340 - lr: 0.0010
Epoch 22/100
44/44 [=====] - 31s 697ms/step - loss: 0.1986 -
accuracy: 0.9307 - lr: 0.0010
Epoch 23/100
44/44 [=====] - 31s 693ms/step - loss: 0.2023 -
accuracy: 0.9287 - lr: 0.0010
Epoch 24/100
44/44 [=====] - 31s 695ms/step - loss: 0.1646 -
accuracy: 0.9465 - lr: 0.0010
Epoch 25/100
44/44 [=====] - 31s 693ms/step - loss: 0.1727 -
accuracy: 0.9407 - lr: 0.0010

```
Epoch 26/100
44/44 [=====] - 32s 719ms/step - loss: 0.1285 -
accuracy: 0.9569 - lr: 0.0010
Epoch 27/100
44/44 [=====] - 31s 703ms/step - loss: 0.1705 -
accuracy: 0.9400 - lr: 0.0010
Epoch 28/100
44/44 [=====] - 31s 709ms/step - loss: 0.1531 -
accuracy: 0.9439 - lr: 0.0010
Epoch 29/100
44/44 [=====] - 32s 711ms/step - loss: 0.1413 -
accuracy: 0.9539 - lr: 0.0010
```

```
[ ]: <keras.callbacks.History at 0x7efc457f6e50>
```

```
[ ]: model.save('flower_cnn.h5')
```

```
#Test
```

```
[ ]: import numpy as np
from tensorflow.keras.preprocessing import image
```

```
[ ]: img = image.load_img('/content/test_image.jpg',target_size=(64,64))
img
```

```
[ ]:
```



```
[ ]: h = image.img_to_array(img)
h
```

```
[ ]: array([[1., 1., 1.],
          [1., 1., 1.],
          [1., 1., 1.],
          ...,
          [1., 1., 1.],
          [1., 1., 1.],
          [1., 1., 1.]],

          [[1., 1., 1.],
          [1., 1., 1.],
          [1., 1., 1.],
          ...,
          [1., 1., 1.],
          [1., 1., 1.],
          [1., 1., 1.]])
```

```

        [1., 1., 1.],
        [1., 1., 1.],
        [1., 1., 1.]],

[[1., 1., 1.],
 [1., 1., 1.],
 [1., 1., 1.],
 ...,
 [1., 1., 1.],
 [1., 1., 1.],
 [1., 1., 1.]],

...,

[[1., 1., 1.],
 [1., 1., 1.],
 [1., 1., 1.],
 ...,
 [1., 1., 1.],
 [1., 1., 1.],
 [1., 1., 1.]],

[[1., 1., 1.],
 [1., 1., 1.],
 [1., 1., 1.],
 ...,
 [1., 1., 1.],
 [1., 1., 1.],
 [1., 1., 1.]],

[[1., 1., 1.],
 [1., 1., 1.],
 [1., 1., 1.],
 ...,
 [1., 1., 1.],
 [1., 1., 1.],
 [1., 1., 1.]]], dtype=float32)

```

```

[ ]: h= np.expand_dims(h,axis= 0)
     h

```

```

[ ]: array([[[[1., 1., 1.],
               [1., 1., 1.],
               [1., 1., 1.],
               ...,
               [1., 1., 1.],
               [1., 1., 1.],

```

```

        [1., 1., 1.]],

[[1., 1., 1.],
 [1., 1., 1.],
 [1., 1., 1.],
 ...,
 [1., 1., 1.],
 [1., 1., 1.],
 [1., 1., 1.]],

[[1., 1., 1.],
 [1., 1., 1.],
 [1., 1., 1.],
 ...,
 [1., 1., 1.],
 [1., 1., 1.],
 [1., 1., 1.]],

...,

[[1., 1., 1.],
 [1., 1., 1.],
 [1., 1., 1.],
 ...,
 [1., 1., 1.],
 [1., 1., 1.],
 [1., 1., 1.]],

[[1., 1., 1.],
 [1., 1., 1.],
 [1., 1., 1.],
 ...,
 [1., 1., 1.],
 [1., 1., 1.],
 [1., 1., 1.]],

[[1., 1., 1.],
 [1., 1., 1.],
 [1., 1., 1.],
 ...,
 [1., 1., 1.],
 [1., 1., 1.],
 [1., 1., 1.]]], dtype=float32)

```

```

[ ]: val = list(xtrain.class_indices.keys())
      val

```

```
[ ]: ['daisy', 'dandelion', 'rose', 'sunflower', 'tulip']
```

```
[ ]: val[np.argmax(model.predict(h))]
```

```
[ ]: 'rose'
```